

ARCOS Group

uc3m | Universidad **Carlos III** de Madrid

Lesson 2 (II) Floating point

Computer Structure
Bachelor in Computer Science and Engineering



Contents

1. Introduction

1. Motivation and goals
2. Positional (numeral) systems

2. Representations

1. Alphanumeric

1. Characters
2. Strings

2. Numerical

1. Natural and integer
2. Fixed point
3. Floating point (IEEE 754 standard)

Contents

1. Introduction

1. Motivation and goals
2. Positional (numeral) systems

2. Representations

1. Alphanumeric


1. Characters
2. Strings

2. Numerical

1. Natural and integer
2. **Fixed point**
3. Floating point (IEEE 754 standard)


Reminder: **we need...**

- To know possible representations:




- To know the **characteristics** of theses representations:

- ## ► Limitations



- To know **how work** with the selected representation:



More representation necessities...

► How to represent?

- Very large numbers: $30.556.926.000_{(10)}$
- Very small numbers: $0.0000000000529177_{(10)}$
- Fractional numbers: 1.58567

Reminder

Example of failure...

- ▶ **Ariane 5 explosion (first flight)**
 - ▶ Sent by ESA in June 1996
 - ▶ Cost of development:
10 years and 7 billion dollars
 - ▶ Exploded 40 seconds after launch, at 3700 meters altitude.
 - ▶ Failure due to total loss of altitude information:
 - ▶ The inertial reference system software performed the conversion of a 64-bit floating point real value to a 16-bit integer value.
 - ▶ The number to be stored was greater than 32767 (the largest 16-bit signed integer) and a conversion failure and exception occurred.



Fixed point [rationals]


- ▶ The position of the binary point is fixed and the weights associated with the decimal places are used.

- ▶ Example:

$$1001.1010 = 2^4 + 2^0 + 2^{-1} + 2^{-3} = 9.625$$

Fractional values in binary with fixed point

► Example with 6 bits:



- Example:

$$10.1010_{(2)} = 1 \times 2^1 + 1 \times 2^{-1} + 1 \times 2^{-3} = 2.625_{10}$$

- Using this fixed point, the range is:

□ [0 a 3.9375 (almost 4)]

Fractional powers of 2

i	2^{-i}	
0	1.0	1
1	0.5	1/2
2	0.25	1/4
3	0.125	1/8
4	0.0625	1/16
5	0.03125	1/32
6	0.015625	
7	0.0078125	
8	0.00390625	
9	0.001953125	
10	0.0009765625	

Contents

1. Introduction

1. Motivation and goals
2. Positional (numeral) systems

2. Representations

1. Alphanumeric

1. Characters
2. Strings

2. Numerical


1. Natural and integer
2. Fixed point
3. **Floating point (IEEE 754 standard)**

Floating-point numbers

The diagram shows the scientific notation 9.12×10^{25} . The number 9.12 is underlined and labeled 'mantissa' with an upward arrow. The '10' is labeled 'radix (base)' with an upward arrow. The 'x' is a multiplier. The '25' is labeled 'exponent' with an upward arrow.

- ▶ Each number has a mantissa and an **exponent**
- ▶ Scientific notation (in decimal): normalized form
 - ▶ Only one digit different to 0 on the left of decimal point
- ▶ The number is adapted to the **order of magnitude** of the value to be represented, by translating the *decimal point* by using the exponent

Scientific notation in binary



- ▶ Normalized form:
One 1 (only one digit) in the left of the binary point
 - ▶ Normalized: 1.0001×2^{-9} ,
 - ▶ Not normalized: 0.0011×2^{-8} , 10.0×2^{-10}

IEEE 754 Floating Point Standard

[rationals]



- ▶ Floating point standard used in most computers.
- ▶ **Characteristics** (unless special cases):
 - ▶ Exponent: excess-k with bias $k = 2^{\text{num_bits_in_exponent}} - 1$
 - ▶ Mantissa: sign-magnitude, normalized, with implicit bit
- ▶ Different **formats**:
 - ▶ **Single precision**: 32 bits (sign: 1, exponent: 8, mantissa: 23 and bias: 127)
 - ▶ **Double precision**: 64 bits (sign: 1, exponent: 11, mantissa: 52 and bias: 1023)
 - ▶ **Quad-precision**: 128 bits (sign: 1, exponent: 15, mantissa: 112 and bias: 16383)

Normalization and implicit bit

► Normalization

In order to normalize the mantissa, the exponent is adjusted to have a most significant bit of value 1

► Example: 100100000000000000000000 $\times 2^3$ (already normalized)

► Example: 000100000000010101 $\times 2^3$ (is not)


100000000010101000 $\times 2^0$ (now it is)

► Implicit bit

Once normalized, since the most significant bit is 1, it is **not** stored to leave space for one more bit (increases accuracy).

► This makes it possible to represent mantissa with one bit more

IEEE Standard 754 (single precision)



- ▶ The value is computed (unless special cases) as:

$$\mathbf{N = (-1)^S \times 2^{E-127} \times 1.M}$$

where:

$S = 0$ for positive numbers, $S = 1$ for negative numbers

$0 < E < 255$ ($E=0$ y $E=255$ are special cases)

$000000000000000000000000 \leq M \leq 111111111111111111111111$

IEEE Standard 754 (single precision)

[rationals]

► Special cases:

$$(-1)^s \times 0.\text{mantissa} \times 2^{-126}$$

Exponent	Mantissa	Special value
0 (0000 0000)	0	+/- 0 (depends on sign)
0 (0000 0000)	$\neq 0$	Number NOT normalized
255 (1111 1111)	$\neq 0$	NaN (0/0, sqrt(-4),)
255 (1111 1111)	0	+/- infinite (depends on sign)
1-254	Any	Normalized number (no special)

$$(-1)^s \times 1.\text{mantissa} \times 2^{\text{exponent}-127}$$

Examples

S	E	M	N
1	00000000	000000000000000000000000	-0 (Exception 0) E=0 y M=0.
1	01111111	000000000000000000000000	$-2^0 \times 1.0_2 = -1$
0	10000001	111000000000000000000000	$+2^2 \times 1.111_2 = +2^2 \times (2^0 + 2^{-1} + 2^{-2} + 2^{-3}) = +7.5$
0	11111111	000000000000000000000000	∞ (Exception ∞) E=255 y M=0
0	11111111	100000000000000000000001	NaN (Not a Number) E=255 y M \neq 0.

Example

- a) Calculate the value in decimal associated to this number
0 10000011 110000000000000000000000
represented in IEEE 754 single precision

Example (solution)

- a) Calculate the value in decimal associated to this number
0 10000011 110000000000000000000000
represented in IEEE 754 single precision

- a) Sign bit: $0 \Rightarrow (-1)^0 = +1$
- b) Exponent: $10000011_2 = 131_{10} \Rightarrow E - 127 = 131 - 127 = 4$
- c) Mantissa: $110000000000000000000000 \Rightarrow 1 \times 2^{-1} + 1 \times 2^{-2} = 0.75$

The decimal value is $+1 \times 2^4 \times 1.75 = +28$

Exercise

- b) Represent the number -9 using IEEE 754 single precision

Exercise (Solution)

b) Represent the number -9 using IEEE 754 single precision

$$-9_{10} = -1001_2 = -1001_2 \times 2^0 = -1.001_2 \times 2^3 \text{ (normalized mantissa)}$$

a) Sign: negative $\Rightarrow S=1$

b) Exponent: $3+127 \text{ (bias)} = 130 \Rightarrow 10000010$

c) Mantissa: $1.001 \text{ (impl. bit)} \Rightarrow 001000000000000000000000$

-9 is represented by $1 \ 10000010 \ 001000000000000000000000$

IEEE Standard 754 (single precision) [rationals]

- ▶ Range of representable magnitudes (regardless of sign):

- ▶ Smallest normalized:

- $2^{-127} \times 1.000000000000000000000000_2$

- ▶ Largest normalized:

- $2^{254-127} \times 1.111111111111111111111111_2$

- ▶ Smallest not normalized :

- $2^{-126} \times 0.000000000000000000000001_2$

- ▶ Largest not normalized :

- $2^{-126} \times 0.111111111111111111111111_2$

Exponent	Mantissa	Special value
0	$\neq 0$	Not normalized
1-254	any	Normalized

$$(-1)^s * 0.\text{mantisa} * 2^{-126}$$

$$(-1)^s * 1.\text{mantisa} * 2^{\text{exponente}-127}$$

IEEE Standard 754 (single precision)

[rationals]

► Range of representable magnitudes (regardless of sign):

► Smallest normalized:

$$2^{-127} \times 1.000000000000000000000000_2 = 2^{-126}$$

► Largest normalized:

$$2^{254-127} \times 1.111111111111111111111111_2 = 2^{127} \times (2 - 2^{-23}) = 2^{128} \times (1 - 2^{-24})$$

► Smallest not normalized :

$$2^{-126} \times 0.000000000000000000000001_2 = 2^{-149}$$

► Largest not normalized :

$$2^{-126} \times 0.111111111111111111111111_2 = 2^{-126} \times (1 - 2^{-23})$$

Tip:


$$\begin{array}{rcl} & 1.111111111111111111111111_2 & = X \\ + & 0.000000000000000000000001_2 & = 2^{-23} \\ \hline \end{array}$$

$$10.000000000000000000000000_2 = 2$$

$$X = 2 - 2^{-23}$$

IEEE Standard 754 (single precision) [rationals]

- ▶ Range of representable magnitudes (regardless of sign):
 - ▶ Smallest normalized:
 $2^{-127} \times 1.000000000000000000000000_2 = 2^{-126} = 2^{-127} \times 0.5$
 - ▶ Largest normalized:
 $2^{254-127} \times 1.111111111111111111111111_2 = 2^{127} \times (2 - 2^{-23}) = 2^{128} \times (1 - 2^{-24})$
 - ▶ Smallest not normalized :
 $2^{-126} \times 0.000000000000000000000001_2 = 2^{-149}$
 - ▶ Largest not normalized :
 $2^{-126} \times 0.111111111111111111111111_2 = 2^{-126} \times (1 - 2^{-23})$



Exercise

- ▶ How many *floats* (single precision floating point numbers) are between 1 and 2 (not included)?


- ▶ How many *float* (single precision floating point numbers) are between 2 and 3 (not included)?

Exercise (Solution)


- ▶ How many *floats* (single precision floating point numbers) are between 1 and 2 (not included)?
 - ▶ $1 = 1.000000000000000000000000 \times 2^0$
 - ▶ $2 = 1.000000000000000000000000 \times 2^1$
 - ▶ Between 1 and 2 there are 2^{23} numbers
- ▶ How many *float* (single precision floating point numbers) are between 2 and 3 (not included)?
 - ▶ $2 = 1.000000000000000000000000 \times 2^1$
 - ▶ $3 = 1.100000000000000000000000 \times 2^1$
 - ▶ Between 2 and 3 there are 2^{22} numbers

Discrete representation


- ▶ Variable resolution:
denser near zero, less towards infinity



Representable numbers



(a) Two's complement integers



(b) Floating point numbers

Example 1

inaccuracy

0.4 →

0	0111101	10011001100110011001101
---	---------	-------------------------



3.9999998 e-1

0.1 →

0	01111011	10011001100110011001100
---	----------	-------------------------



9.9999994 e-2

Example 2

inaccuracy

- ▶ How does C performs a division?

t2.c

```
#include <stdio.h>

int main ( )
{
    float a ;

    a = 3.0/7.0 ;
    if (a == 3.0/7.0)
        printf("Equal\n") ;
    else printf("Not equal\n") ;
    return (0) ;
}
```

Example 2

inaccuracy

- ▶ How does C performs a division?

t2.c

```
#include <stdio.h>

int main ( )
{
    float a ;

    a = 3.0/7.0 ;
    if (a == 3.0/7.0)
        printf("Equal\n") ;
    else printf("Not equal\n") ;
    return (0) ;
}
```

```
$ gcc -o t2 t2.c
$ ./t2
Not equal
```

Example 2

inaccuracy

- ▶ How does C performs a division?

t2.c

```
#include <stdio.h>
```

```
int main ( )
```

```
{
```

```
    float a ;
```

float

```
    a = 3.0/7.0 ;
```

```
    if (a == 3.0/7.0)
```

double

```
        printf("Equal\n") ;
```

```
    else printf("Not equal\n") ;
```

```
    return (0) ;
```

```
}
```

```
$ gcc -o t2 t2.c
```

```
$ ./t2
```

```
Not equal
```


Example 3

inaccuracy

- ▶ The associative property is not always satisfied
 $a + (b + c) = (a + b) + c$?

t1.c

```
#include <stdio.h>

int main ( )
{
    float x, y, z ;

    x = 10e30;  y = -10e30;  z = 1;
    printf("(x+y)+z = %f\n", (x+y)+z) ;
    printf("x+(y+z) = %f\n", x+(y+z)) ;

    return (0) ;
}
```

Example 3

inaccuracy

- ▶ The associative property is not always satisfied
 $a + (b + c) = (a + b) + c$?

t1.c

```
#include <stdio.h>

int main ( )
{
    float x, y, z ;

    x = 10e30;  y = -10e30;  z = 1;
    printf("(x+y)+z = %f\n", (x+y)+z) ;
    printf("x+(y+z) = %f\n", x+(y+z)) ;

    return (0) ;
}
```

```
$ gcc -o t1 t1.c
```

```
$ ./t1
```

```
(x+y)+z = 1.000000
```

```
x+(y+z) = 0.000000
```

Floating-point is not associative

- ▶ Floating-point is not associative

- ▶ $x = -1.5 \times 10^{38}$, $y = 1.5 \times 10^{38}$, $z = 1.0$

- ▶ $x + (y + z) = -1.5 \times 10^{38} + (1.5 \times 10^{38} + 1.0)$
 $= -1.5 \times 10^{38} + (1.5 \times 10^{38}) = 0.0$

- ▶ $(x + y) + z = (-1.5 \times 10^{38} + 1.5 \times 10^{38}) + 1.0$
 $= (0.0) + 1.0 = 1.0$

- ▶ Floating point operations are not associative

- ▶ Results are approximated

- ▶ 1.5×10^{38} is so much larger than 1.0

- ▶ $1.5 \times 10^{38} + 1.0$ in floating point representation is still 1.5×10^{38}

Example

int → float → int

```
if (i == (int) ((float) i)) {  
    printf("true");  
}
```

- ▶ **Not** always prints "true"
- ▶ Most integer values (specially larger ones) don't have an exact floating point representation
- ▶ What about double?

Example

- ▶ The number 133000405 in binary is:
 - ▶ 111111011010110110011010101 (27 bits)
- ▶ 111111011010110110011010101 $\times 2^0$
- ▶ When is normalized:
 - ▶ 1.11111011010110110011010101 $\times 2^{26}$
 - ▶ $S = 0$ (positive)
 - ▶ $e = 26 \rightarrow E = 26 + 127 = 153$
 - ▶ $M = 11111011010110110011010$ (last 3 bits are lost)
- ▶ The normalized number stored is:
 - ▶ 1.11111011010110110011010 $\times 2^{26} =$
 - ▶ 111111011010110110011010 $\times 2^3 = 133000400$

Example

float → **int** → **float**

```
if (f == (float)((int) f)) {  
    printf("true");  
}
```

- ▶ Not always true
- ▶ Numbers with decimals do not have integer representation

Rounding

- ▶ Rounding removes less significant digits from a number to obtain an approximate value.
- ▶ **Types** of rounding:
 - ▶ Round **to + ∞**
 - ▶ Round it “up”: $2.001 \rightarrow 3$, $-2.001 \rightarrow -2$
 - ▶ Round **to - ∞**
 - ▶ Round it “down”: $1.999 \rightarrow 1$, $-1.999 \rightarrow -2$
 - ▶ **Truncate**
 - ▶ Discard last bits: $1.299 \rightarrow 1.2$
 - ▶ Round **to nearest (ties to even)**
 - ▶ $2.4 \rightarrow 2$, $2.6 \rightarrow 3$, $-1.4 \rightarrow -1$
 - ▶ If number falls midway then it is rounded to the nearest value with an even least significant digit ($+23.5 \rightarrow +24 \leftarrow +24.5$; $-23.5 \rightarrow -24 \leftarrow -24.5$)

Rounding

- ▶ Rounding means losing accuracy.
- ▶ Rounding occurs:
 - ▶ When moving to a representation with fewer representables:
 - ▶ E.g.: A value from double to single precision
 - ▶ E.g.: A floating point value to integer
 - ▶ When performing arithmetic operations:
 - ▶ E.g.: After adding two floating-point numbers (using guard bits)

Guard bits

- ▶ **Guard digits** are used to improve accuracy:
 - ▶ FP hardware internally includes additional bits for operations
 - ▶ After operation, guard bits are eliminated: rounding
- ▶ Example: $2.65 \times 10^0 + 2.34 \times 10^2$

	WITHOUT guard bits	WITH guard bits
1.- equalize exponents	0.02×10^2 $+ 2.34 \times 10^2$	$0.02\textcolor{blue}{65} \times 10^2$ $+ 2.34\textcolor{blue}{00} \times 10^2$
2.- add	2.36×10^2	$2.36\textcolor{blue}{65} \times 10^2$
3.- round	$2.3\textcolor{red}{6} \times 10^2$	$2.3\textcolor{red}{7} \times 10^2$

Floating point operations

- ▶ **Add**

- ▶ **Subtract**


1. Check zero values.
2. Equalize exponents (shift smaller number to the right).
3. Add/subtract mantissa.
4. Normalize the result.

- ▶ **Multiply**

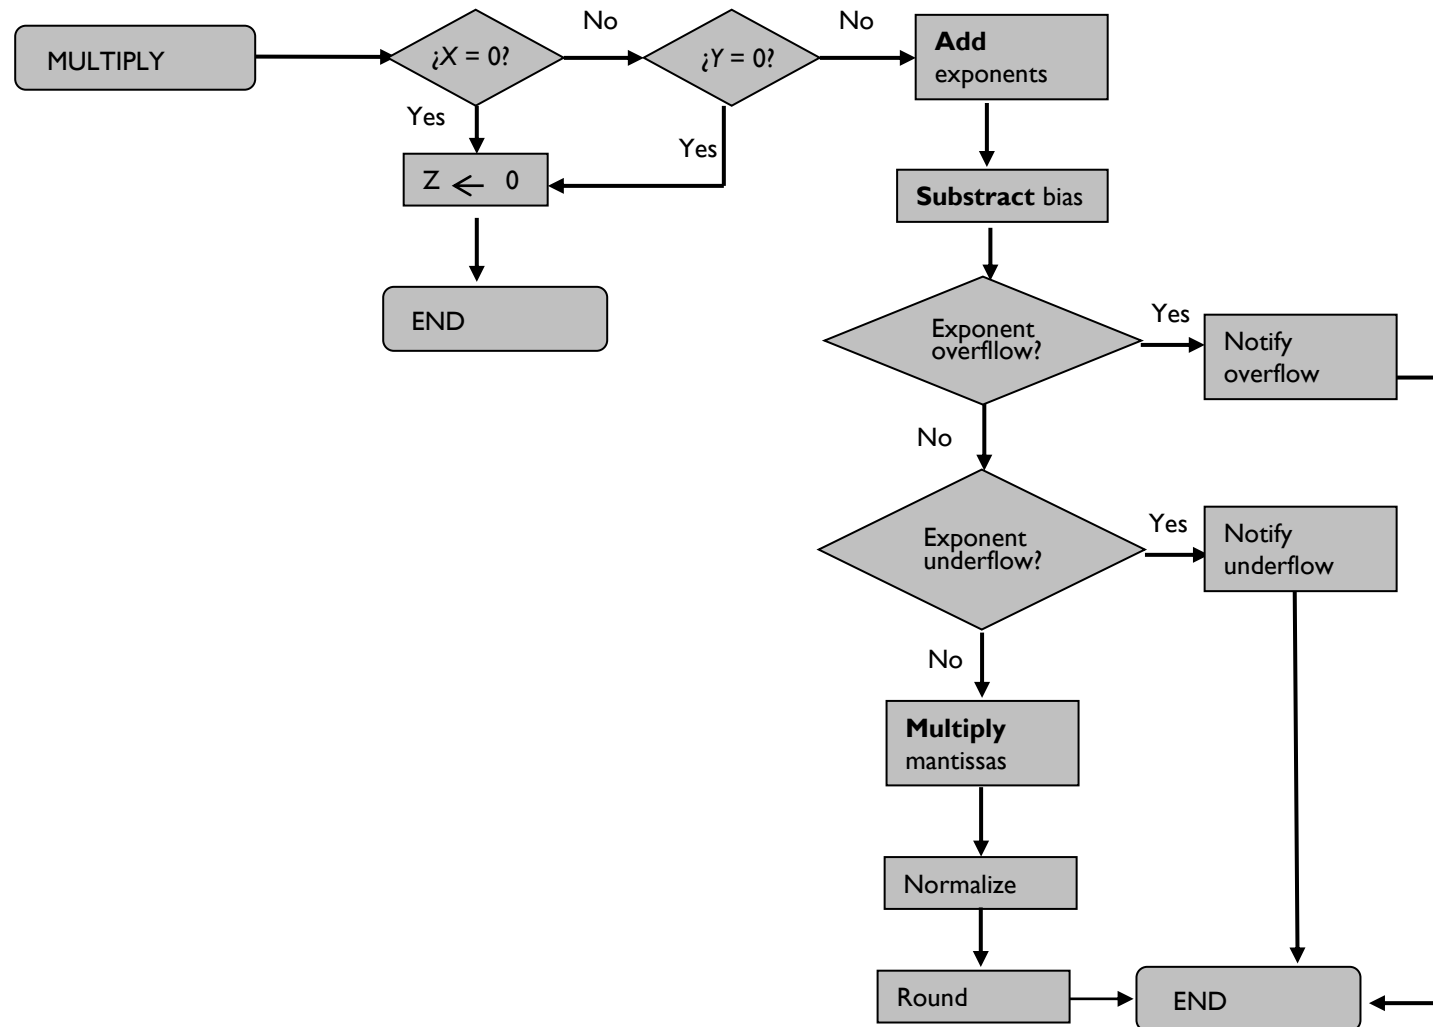
- ▶ **Divide**

1. Check zero values.
2. Add/subtract exponents.
3. Multiply/divide mantissa (taking into account the sign).
4. Normalize the result.
5. Rounding the result.

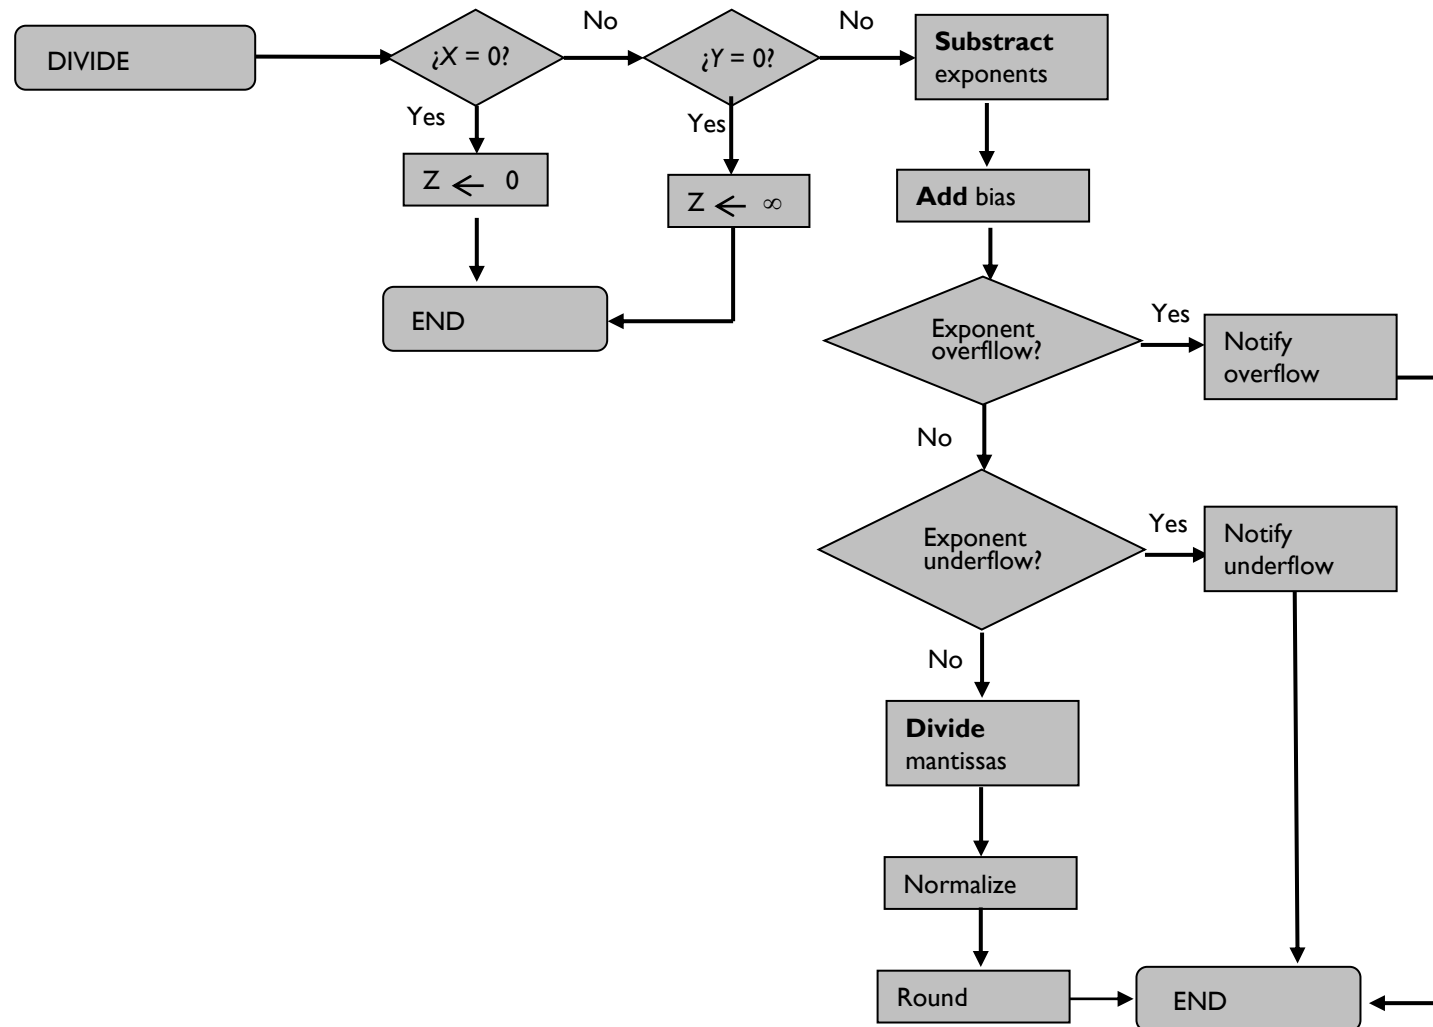
Additions and subtractions: $Z=X+Y$ y $Z=X-Y$



Multiplication: $Z = X * Y$



Division: $Z = X/Y$



Exercise

- ▶ Using the IEEE 754 format, add 7.5 and 1.5 step by step.

Solution

To binary

1) $7.5 + 1.5 =$

2) $1.111 * 2^2 + 1.1 * 2^0 =$

Equalize
exponents

3) $1.111 * 2^2 + 0.011 * 2^2 =$

4) $10.010 * 2^2 =$

Add

5) $1.0010 * 2^3$

Adjust
exponents

Solution

► Representation of the numbers


7.5 → 0 10000001 111000000000000000000000

+ **1.5** → 0 01111111 100000000000000000000000

- **7.5** = **111.1** × 2⁰ = **1.111** × 2²
Sign = 0 (positive)
Exponent = 2 → exponent to store = 2 + 127 = 129 = 10000001
Mantissa = 1.111 → mantissa to store = 1110000 ... 0000
- **1.5** = **1.1** × 2⁰
Sign = 0 (positive)
Exponent = 0 → exponent to store = 0 + 127 = 127 = 01111111
Mantissa = 1.1 → mantissa to store = 1000000 ... 0000


Solution

- Splitting exponents and mantissas, and adding implicit bit




Solution

► Equalize exponents



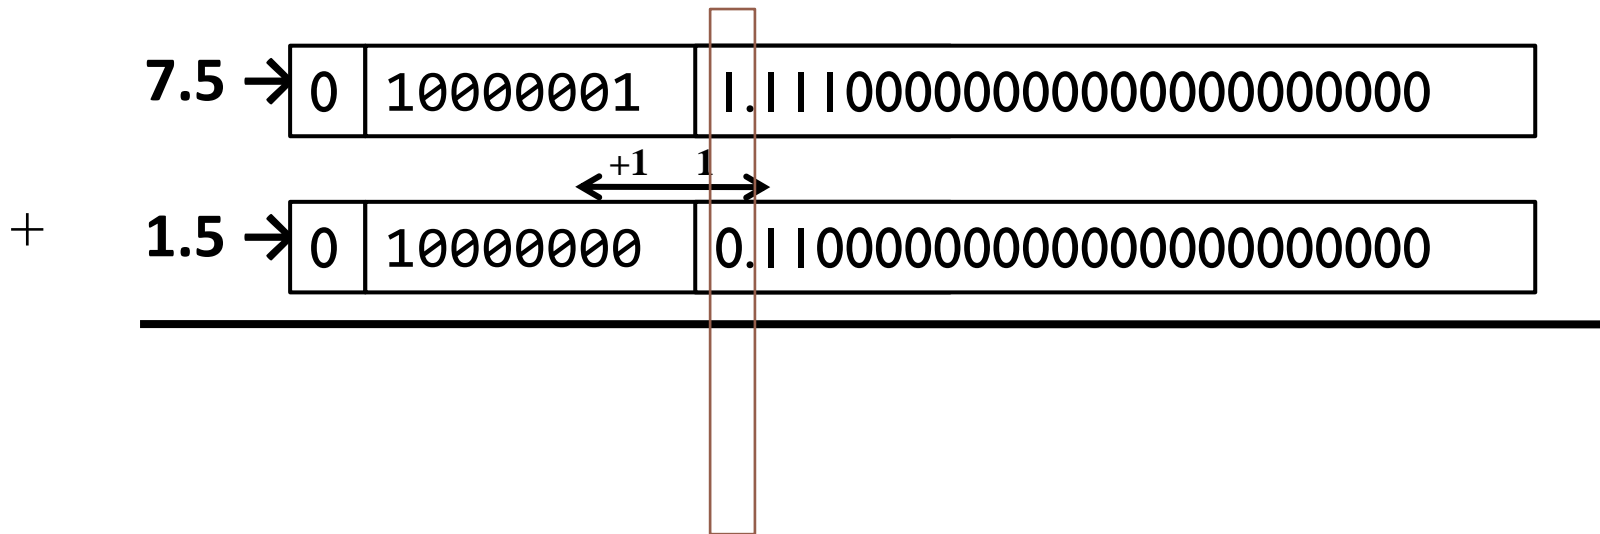
Solution

► Equalize exponents




Solution

► Equalize exponents




Solution

► Equalize exponents




Solution

► Add mantissas



Solution

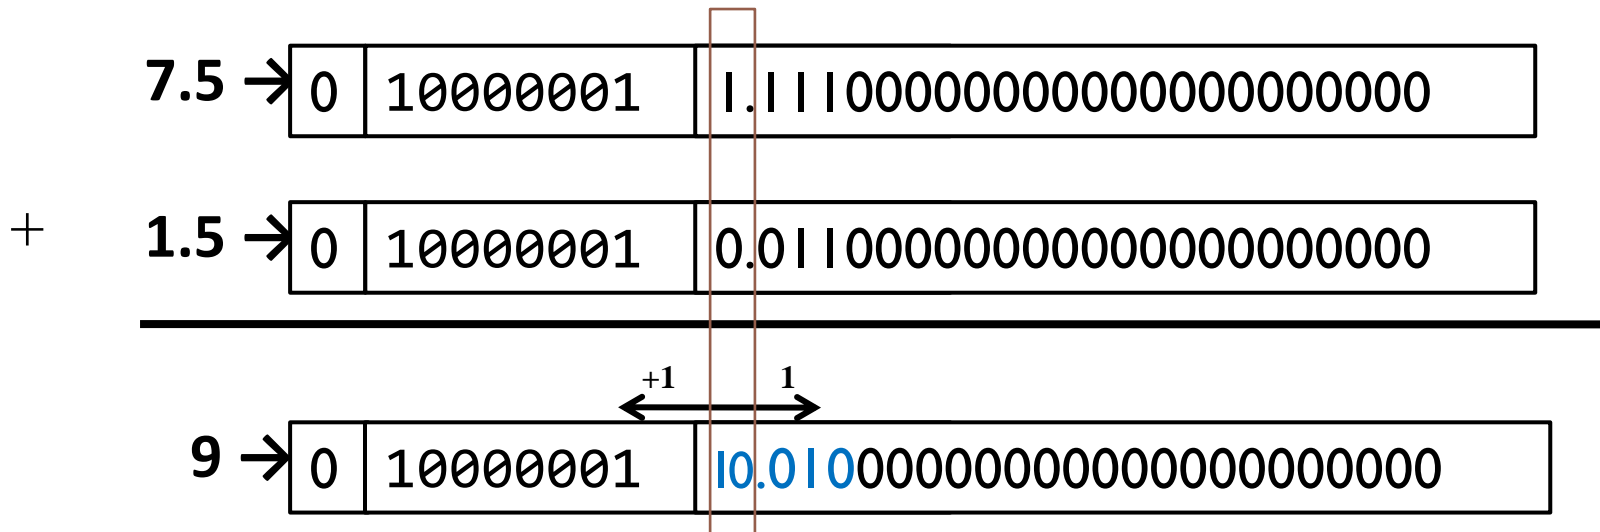
► **Normalize result...**



There is carry,
non-normalized mantissa

Solution

► Normalize result...



There is carry,
non-normalized mantissa

Solution

+

7.5 →	0	10000001	1.11100000000000000000000000000000
1.5 →	0	10000001	0.01100000000000000000000000000000
<hr/>			
9 →	0	10000010	1.00100000000000000000000000000000

Solution

- ▶ Eliminate the implicit bit and store the result

9 → 0 10000010 001000000000000000000000

Exercise

- ▶ Using the IEEE 754 format, compute $9 - 7.5$ step by step.

Solution

► Representation of the numbers

9 → 0 10000010 001000000000000000000000

- 7.5 → 1 10000001 111000000000000000000000

- $-7.5 = 111.1 \times 2^0 = 1.111 \times 2^2$

Sign = 1 (negative)

Exponent = 2 → exponent to store = 2 + 127 = 129 = 10000001

Mantissa = 1.111 → mantissa to store = 1110000 ... 0000

- $9 = 1001.0 \times 2^0 = 1.0010 \times 2^3$

Sign = 0 (positive)

Exponent = 3 → exponent to store = 3 + 127 = 130 = 1000010

Mantissa = 1.001 → mantissa to store = 0010000 ... 0000

Solution

- ▶ Splitting exponents and mantissas, and adding implicit bit

9 →

0	10000010	1.001000000000000000000000
---	----------	----------------------------

- 7.5 →


1	10000001	1.111000000000000000000000
---	----------	----------------------------

adding implicit bit


Solution

► Equalize exponents

9 →




- 7.5 →




Solution

► Equalize exponents



Solution

► Equalize exponents



Solution


► Subtract

$$\begin{array}{r} 9 \rightarrow \begin{array}{|c|c|c|} \hline 0 & 10000010 & 1.001000000000000000000000 \\ \hline \end{array} \\ - 7.5 \rightarrow \begin{array}{|c|c|c|} \hline 1 & 10000010 & 0.111100000000000000000000 \\ \hline \end{array} \\ \hline 1.5 \rightarrow \begin{array}{|c|c|c|} \hline 0 & 10000010 & 0.001100000000000000000000 \\ \hline \end{array} \end{array}$$


Solution

► Normalize result...


9 →



- 7.5 →



1.5 →



Solution

► Normalize result...

9 →

0	10000010	1.001000000000000000000000
---	----------	----------------------------

- 7.5 →

1	10000010	0.111100000000000000000000
---	----------	----------------------------

1.5 →


0	10000001	0.011000000000000000000000
---	----------	----------------------------

-1 1
→ ←


Solution

► Normalize result...


9 →



- 7.5 →




1.5 →




Solution

► Normalize result...


9 →



- 7.5 →



1.5 →



mantissa already normalized

Solution

- ▶ Eliminate the implicit bit and store the result

1.5 → 0 01111111 100000000000000000000000

Exercise

- ▶ Using the IEEE 754 format, multiply 7.5 and 1.5 step by step.

Solution

summary

$$\begin{aligned} 7.5 \times 1.5 &= (1.111_2 \times 2^2) \times (1.1_2 \times 2^0) \\ &= (1.111_2 \times 1.1_2) \times 2^{(2+0)} \\ &= (10.1101_2) \times 2^2 \\ &= (1.01101_2) \times 2^3 \\ &= 11.25 \end{aligned}$$

Solution

► Representation of the numbers

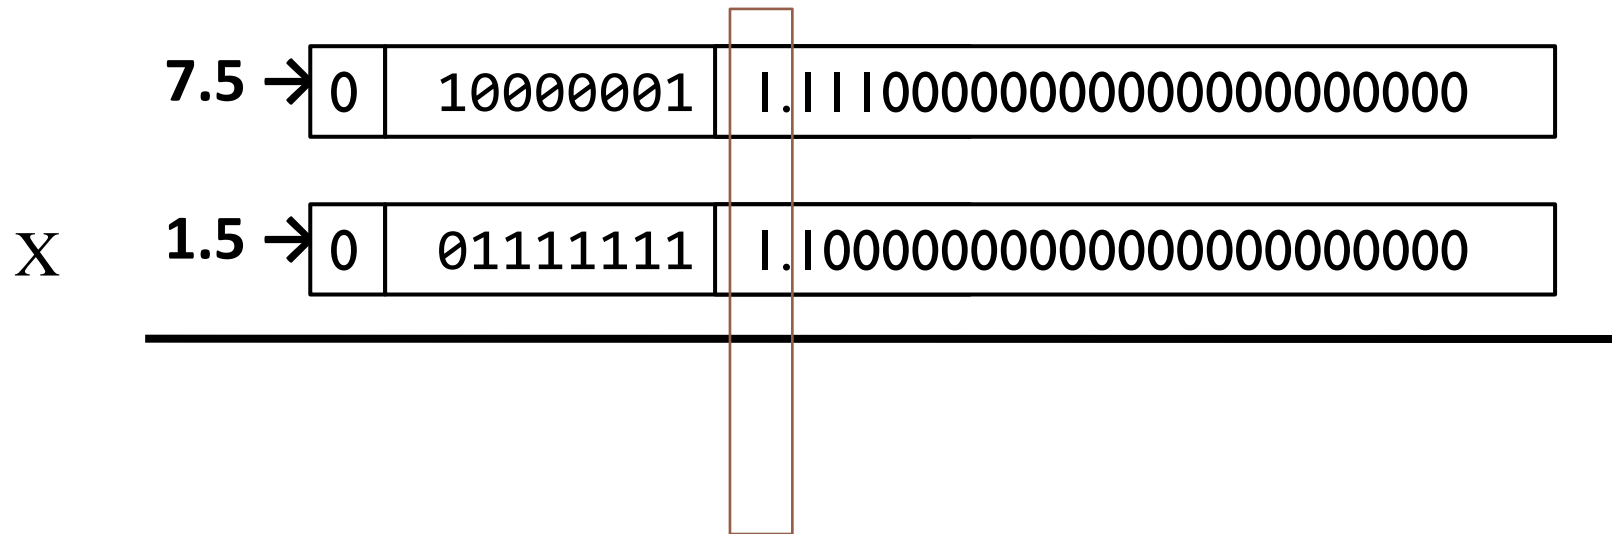
7.5 → 0 10000001 111000000000000000000000

1.5 → 0 01111111 110000000000000000000000

- $7.5 = 111.1 \times 2^0 = 1.111 \times 2^2$
Sign = 0 (positive)
Exponent = 2 → exponent to store = $2 + 127 = 129 = 10000001$
Mantissa = 1.111 → mantissa to store = 1110000 ... 0000
- $1.5 = 1.1 \times 2^0$
Sign = 0 (positive)
Exponent = 0 → exponent to store = $0 + 127 = 127 = 01111111$
Mantissa = 1.1 → mantissa to store = 1000000 ... 0000

Solution

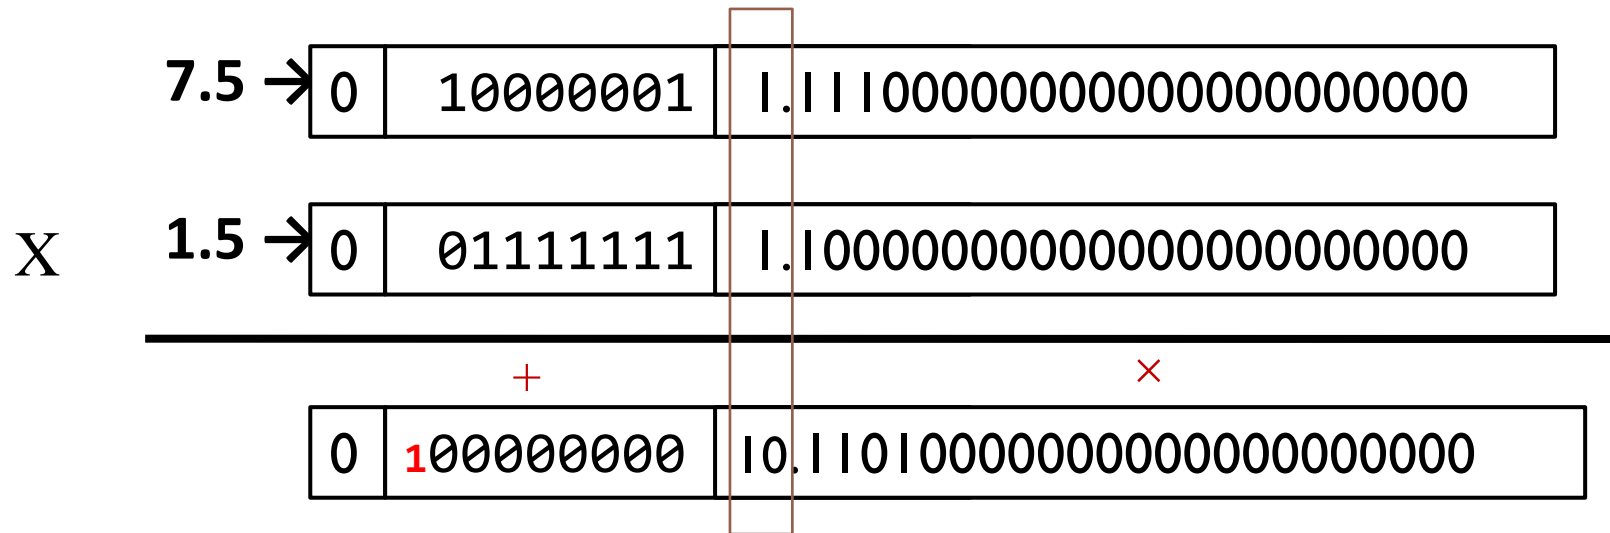
- Splitting exponents and mantissas, and adding implicit bit



The implicit bit is included

Solution

- ▶ Multiply: add exponents and multiply mantissas



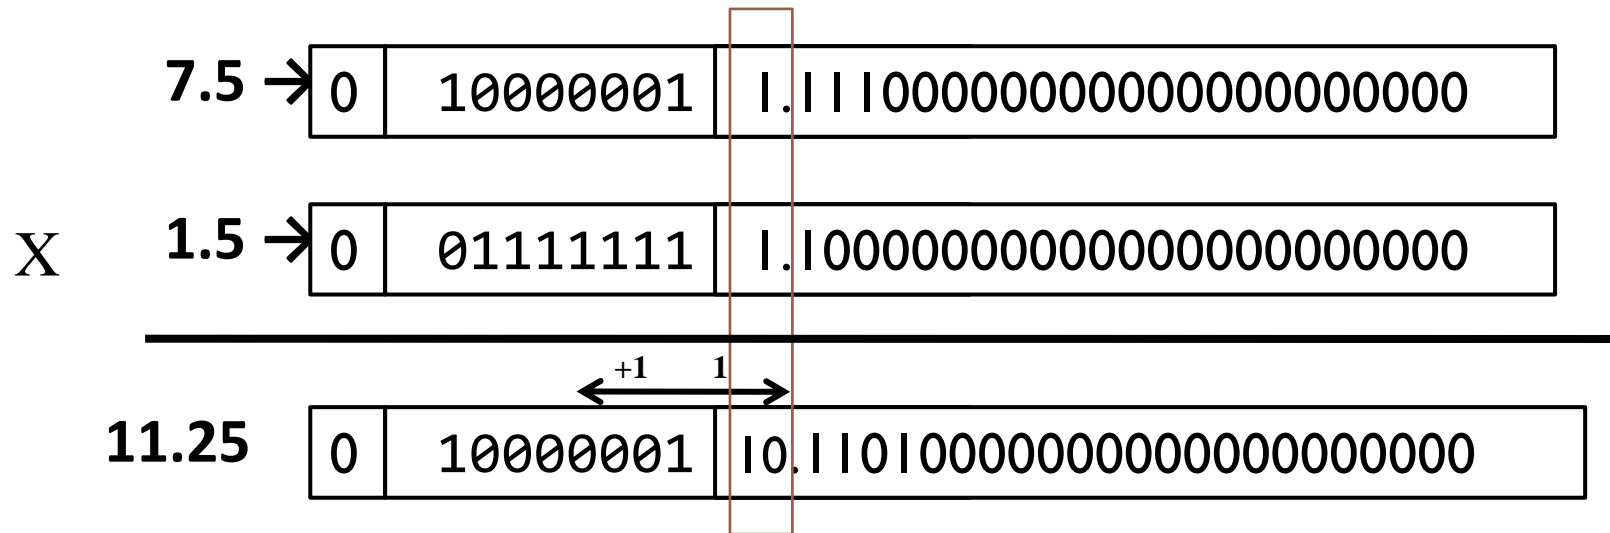
Solution

- ▶ Multiply: remove one bias from exponent (there are two)

$$\begin{array}{r} 7.5 \rightarrow \begin{array}{|c|c|c|} \hline 0 & 10000001 & 1.111000000000000000000000 \\ \hline \end{array} \\ \times 1.5 \rightarrow \begin{array}{|c|c|c|} \hline 0 & 01111111 & 1.100000000000000000000000 \\ \hline \end{array} \\ \hline \begin{array}{|c|c|c|} \hline 0 & 10000000 & 10.110100000000000000000000 \\ \hline \end{array} \\ - \quad \begin{array}{|c|c|c|} \hline 0 & 01111111 & \\ \hline \end{array} \\ \hline \begin{array}{|c|c|c|} \hline 0 & 10000001 & 10.110100000000000000000000 \\ \hline \end{array} \end{array}$$

Solution

- ▶ Multiply: normalize result...



Solution

- ▶ Multiply: normalize result...

	7.5 →	0	10000001	1.111000000000000000000000
X	1.5 →	0	01111111	1.100000000000000000000000
<hr/>				
	11.25	0	10000010	1.011010000000000000000000

Solution

- ▶ Eliminate the implicit bit and store the result

11.25 0 10000010 011010000000000000000000

IEEE 754 Evolution

- ▶ 1985 – IEEE 754
- ▶ 2008 – IEEE 754-2008 (754+854)
- ▶ 2011 – ISO/IEC/IEEE 60559:2011 (754-2008)

Name	Common name	Base	Digits	E min	E max	Notes	Decimal digits	Decimal E max
binary16	Half precision	2	10+1	−14	+15	storage, not basic	3.31	4.51
binary32	Single precision	2	23+1	−126	+127		7.22	38.23
binary64	Double precision	2	52+1	−1022	+1023		15.95	307.95
binary128	Quadruple precision	2	112+1	−16382	+16383		34.02	4931.77
decimal32		10	7	−95	+96	storage, not basic	7	96
decimal64		10	16	−383	+384		16	384
decimal128		10	34	−6143	+6144		34	6144

http://en.wikipedia.org/wiki/IEEE_floating_point

ARCOS Group

uc3m | Universidad **Carlos III** de Madrid

Lesson 2 (II) Floating point

Computer Structure
Bachelor in Computer Science and Engineering



How many not normalized numbers different to zero can be represented?

$$(s) \times 0.\text{mantissa} \times 2^{-126}$$

Exponent	Mantissa	Special value
0 (0000 0000)	No cero	Number not normalized

How many not normalized numbers different to zero can be represented?

Exponent	Mantissa	Special value
0 (0000 0000)	No zero	Number not normalized

$(s) \times 0.\text{mantissa} \times 2^{-126}$

► Solution:

- 23 bits for mantissa (different to 0)

$$2^{23} - 1$$

Example

- ▶ What is the binary and decimal value of the following number represented in the IEEE 754 standard?
3FE00000

Solution

► Binary value:

3	F	E	0	0	0	0	0
0011	1111	1110	0000	0000	0000	0000	0000

► In decimal:

0011 1111 1110 0000 0000 0000 0000 0000

- Sign: 0
- Exponent: 01111111 $\Rightarrow 127-127 = 0$
- Mantissa: 1.110000000000000000000000 $\Rightarrow 1+0.5+0.25 = 1.75$

Then, the value is $+1 \times 1.75 \times 2^0 = 1.75$