

ARCOS Group

uc3m | Universidad **Carlos III** de Madrid

Lesson 4 (III) The processor

Computer Structure
Bachelor in Computer Science and Engineering



Contents

1. Computer elements
2. Processor organization
3. Control unit
4. Execution of instructions
5. Control unit design
6. Execution modes
7. Interrupts
8. Computer startup
9. Performance and parallelism

Contents


1. Computer elements
2. Processor organization
3. Control unit
4. Execution of instructions
5. Control unit design
6. Execution modes
7. Interrupts
8. Computer startup
9. Performance and parallelism

Modes of execution

- ▶ It is indicated by a bit in the status register (U)
- ▶ At least 2 modes:
 - ▶ **User Mode**
 - ▶ The processor cannot **execute privileged instructions** (e.g.: I/O instructions, interrupt enable instructions, ...)
 - ▶ If a user process executes a privileged instruction, an interruption (exception) occurs
 - ▶ **Kernel Mode**
 - ▶ Reserved to the operating system
 - ▶ The processor can execute the entire repertoire of instructions

Activation of the status register

SELEC operation:



if (SelP1 = 1 AND SelP0 == 1)
Output = $C'V'N'Z'I'U$


if (SelP1 == 1 AND SelP0 == 0)
Output = $C'V'N'Z'I'U$

if (SelP1 == 0 AND SelP0 == 1)
Output = $C'V'N'Z'I'U'$

Contents

1. Computer elements
2. Processor organization
3. Control unit
4. Execution of instructions
5. Control unit design
6. Execution modes
7. **Interrupts**
8. Computer startup
9. Performance and parallelism

Interrupts




- ▶ Condition detected by the Control Unit that breaks the normal execution sequence:
 - ▶ The current program is stopped.
 - ▶ The execution is transferred to another program that attend the interruption (Interrupt Service Routine a.k.a. ISR)
 - ▶ When the ISR ends, the execution of the interrupted program is resumed.
- ▶ Example of causes:
 - ▶ When a peripheral requests the attention of the processor,
 - ▶ When an error occurs in the execution of the instruction, Etc.

Classification of interruptions

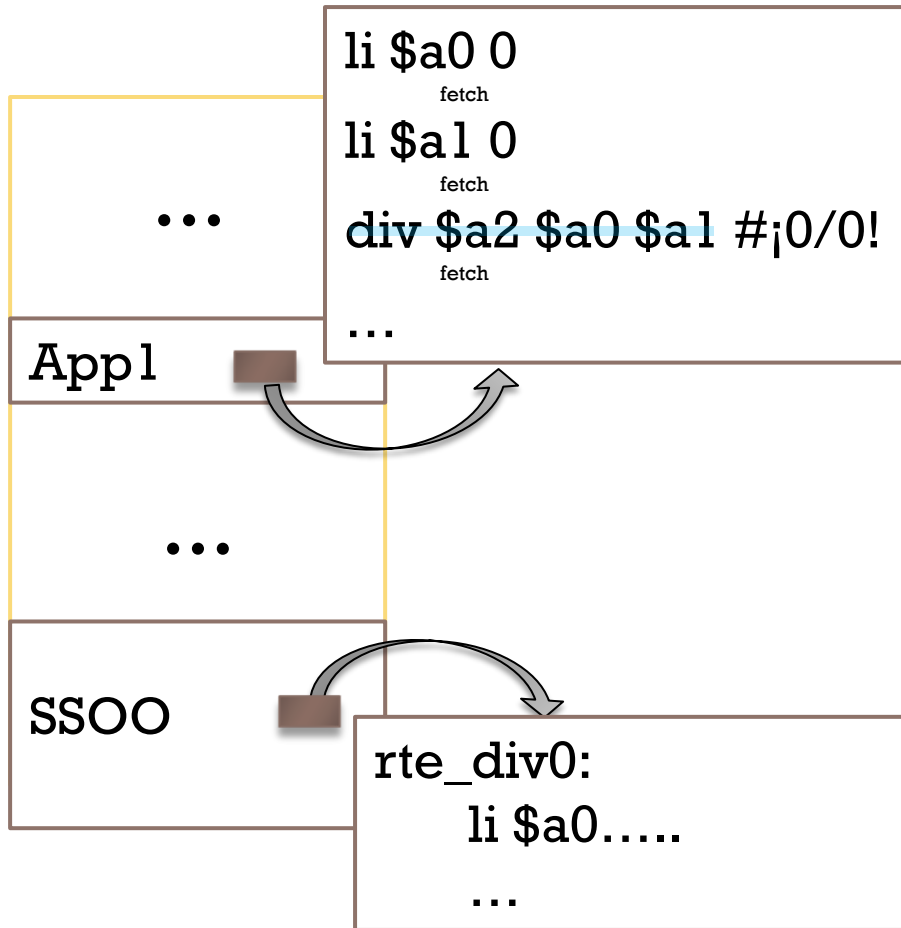
- ▶ **Synchronous** hardware exceptions
 - ▶ When an **error** occurs **in the execution of the instruction**: Division by zero, access to an illegal memory position, illegal instruction, etc.
- ▶ **Asynchronous** hardware exceptions
 - ▶ Faults or **errors** in hardware **not related to current instruction**: printer without paper, power failure, etc.
- ▶ **External interruptions**
 - ▶ When a peripheral (or CPU) requests the attention of the CPU: Peripherals, clock interruption
- ▶ **System calls**
 - ▶ Request an operating system service
 - ▶ Special machine instructions that generate an interruption to activate the operating system

Asynchronous Hardware Exceptions and External Interrupts



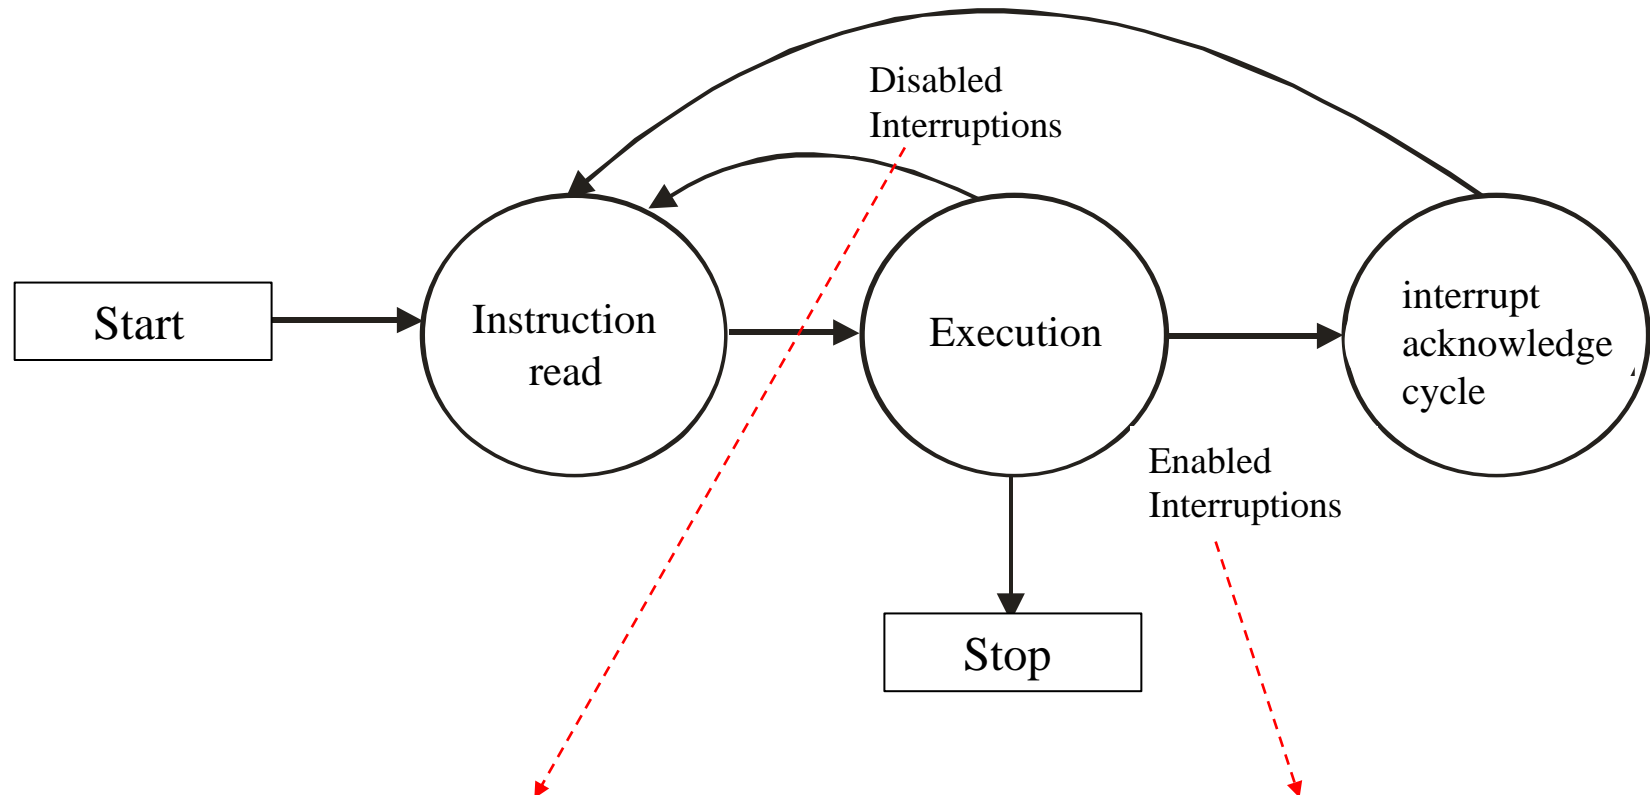
- ▶ They cause an unscheduled sequence break
 - ▶ **Before doing the fetch cycle, first see if there is any pending interruption, and if so...**
 - ▶ ...Bifurcation to subroutine of the O.S. that treats it
- ▶ It then restores the status and returns control to the interrupted program.
- **Asynchronous cause to the execution of the current program**
 - ▶ Peripheral care
 - ▶ Etc.

Synchronous Hardware Exceptions and System Calls



- ▶ They cause an unscheduled sequence break
 - ▶ **Within the microprogram of the ongoing instruction...**
 - ▶ ...Bifurcation to subroutine of the O.S. that treats it
- ▶ It restores the status and returns control to the interrupted program **or ends its execution**
- **Synchronous cause to the execution of the current program**
 - ▶ Division between zero
 - ▶ Etc.


Activation of the status register



It is indicated by a bit located in **the status register (I)**

Activation of the status register

SELEC operation:



if (SelP1 = 1 AND SelP0 == 1)
Output = **C' V' N' Z'** I U

if (SelP1 == 1 AND SelP0 == 0)
Output = C V N Z **I'** U

if (SelP1 == 0 AND SelP0 == 1)
Output = C V N Z I **U'**

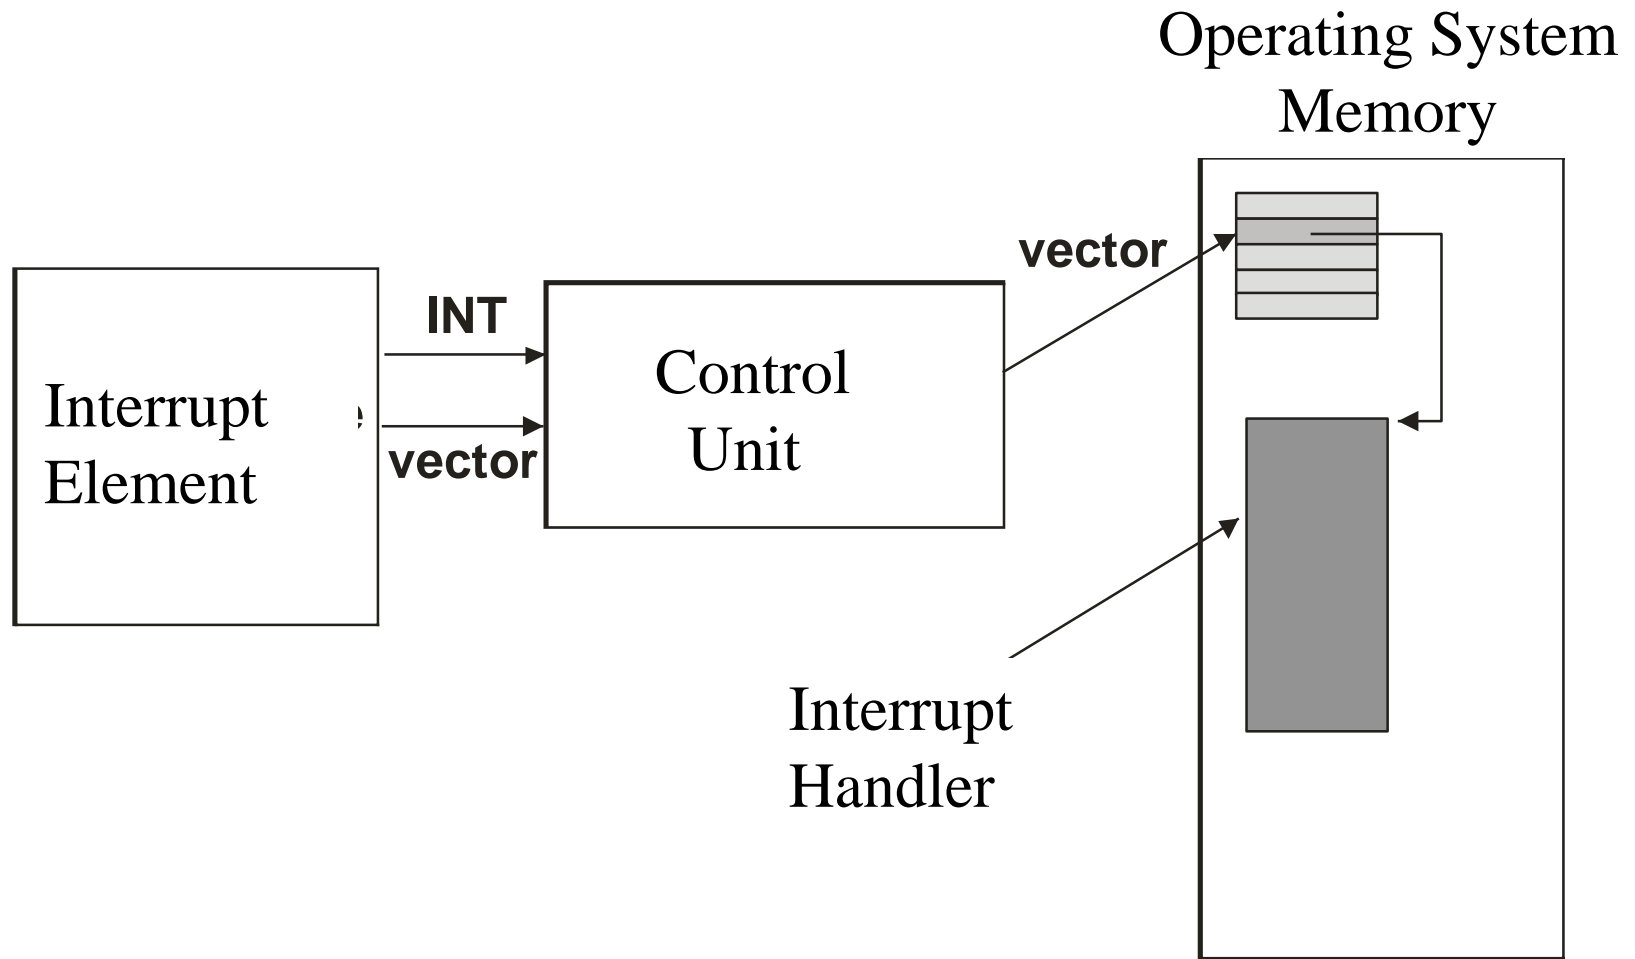
Interrupt Acknowledge Cycle (IAC)

- ▶ It is a microcode before the fetch cycle
 - ▶ It handles the asynchronous interrupts
- ▶ General structure of the IAC:
 1. Checks if an interruption signal is activated.
 2. If it is activated:
 1. Saves the program counter and status register
 - Equivalent to “push pc, push sr”
 2. Switches from user mode to kernel mode
 - Equivalent to “SR.U = 0”
 3. Obtains the address of the Interrupt Service Routine (ISR)
 - Equivalent to “isr_addr = Vector_interrupts[id_interrupt]”
 4. Store the address obtained in the program counter (this way the following instruction will be the first one for the treatment routine)
 - Equivalent to “PC = isr_addr”

Interrupt Service Routine (ISR)

- ▶ It is part of the operating system code
 - ▶ There is one ISR for each interruption that may occur
- ▶ General structure of the ISR:
 1. Saves the rest of the processor registers (if required)
 2. Service the interrupt
 3. Restores processor registers saved in (2)
 4. Executes a special machine instruction: RETI
 - ▶ Resets the status register of the interrupted program (by setting the processor mode back to user mode).
 - ▶ Resets the program counter (so that the next instruction is that of the interrupted program).

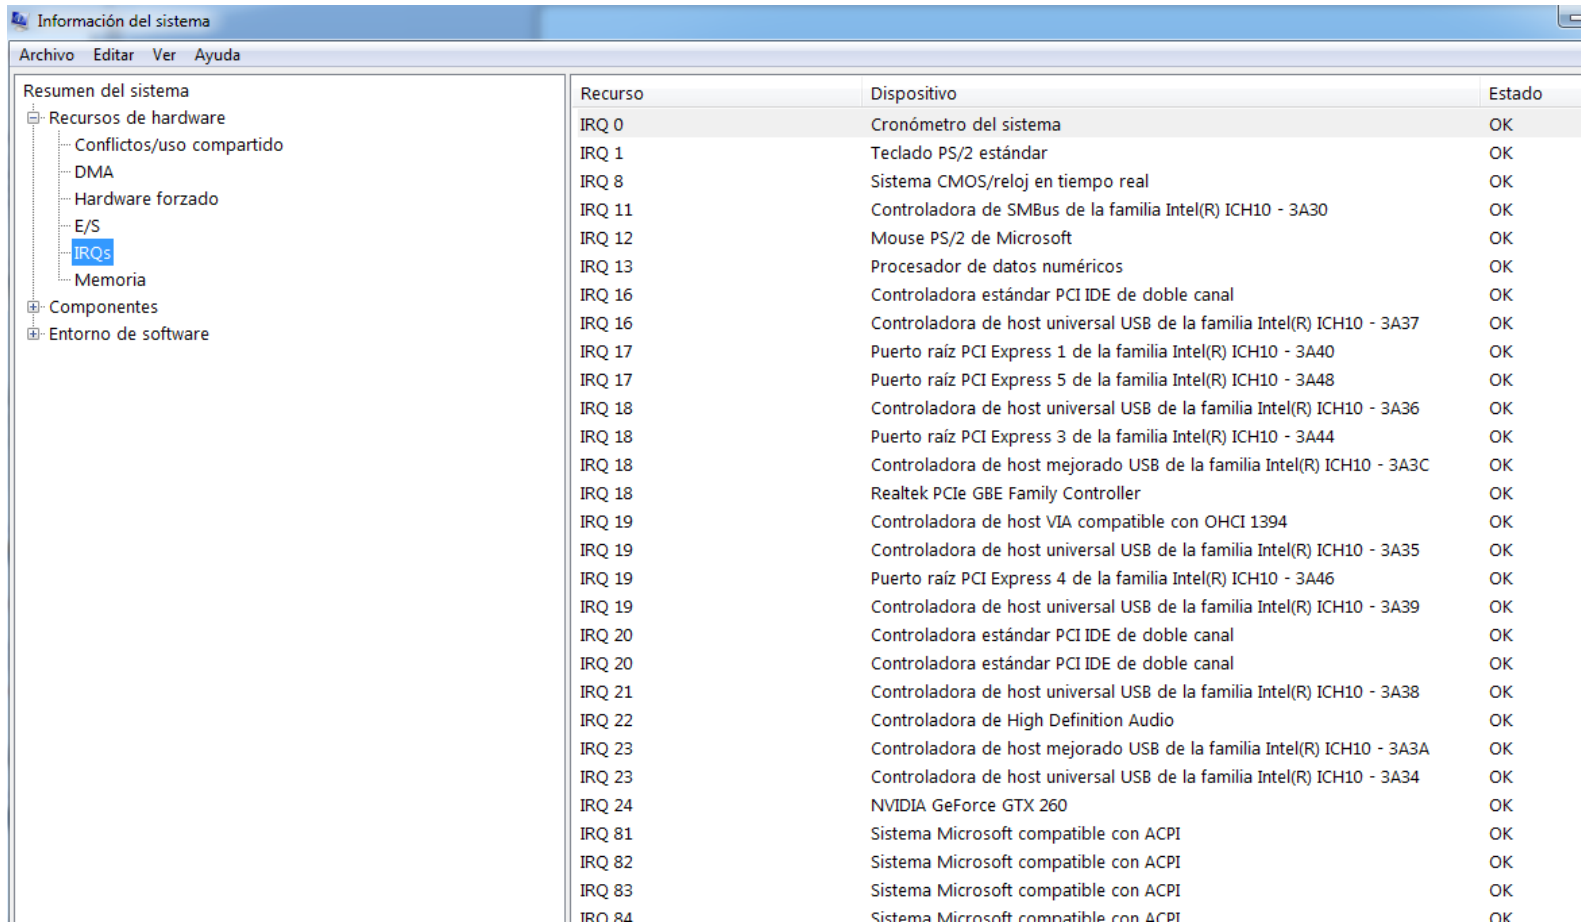
Vector interrupts



Vector interrupts

- ▶ The interrupting element supplies the **interrupt vector**
- ▶ This **vector** is an index in a table containing the address of the interrupt handler routine.
- ▶ The Control Unit reads the content of this entry and loads the value into the PC
- ▶ Each operating system fills this table with the addresses of each of the treatment routines, which are dependent on each operating system.

Interrupts in Windows



The screenshot shows the 'Información del sistema' window in Windows. The left sidebar is expanded to 'Resumen del sistema' > 'Recursos de hardware' > 'IRQs'. The main pane displays a table of system resources and their assigned IRQs.

Recurso	Dispositivo	Estado
IRQ 0	Cronómetro del sistema	OK
IRQ 1	Teclado PS/2 estándar	OK
IRQ 8	Sistema CMOS/reloj en tiempo real	OK
IRQ 11	Controladora de SMBus de la familia Intel(R) ICH10 - 3A30	OK
IRQ 12	Mouse PS/2 de Microsoft	OK
IRQ 13	Procesador de datos numéricos	OK
IRQ 16	Controladora estándar PCI IDE de doble canal	OK
IRQ 16	Controladora de host universal USB de la familia Intel(R) ICH10 - 3A37	OK
IRQ 17	Puerto raíz PCI Express 1 de la familia Intel(R) ICH10 - 3A40	OK
IRQ 17	Puerto raíz PCI Express 5 de la familia Intel(R) ICH10 - 3A48	OK
IRQ 18	Controladora de host universal USB de la familia Intel(R) ICH10 - 3A36	OK
IRQ 18	Puerto raíz PCI Express 3 de la familia Intel(R) ICH10 - 3A44	OK
IRQ 18	Controladora de host mejorado USB de la familia Intel(R) ICH10 - 3A3C	OK
IRQ 18	Realtek PCIe GBE Family Controller	OK
IRQ 19	Controladora de host VIA compatible con OHCI 1394	OK
IRQ 19	Controladora de host universal USB de la familia Intel(R) ICH10 - 3A35	OK
IRQ 19	Puerto raíz PCI Express 4 de la familia Intel(R) ICH10 - 3A46	OK
IRQ 19	Controladora de host universal USB de la familia Intel(R) ICH10 - 3A39	OK
IRQ 20	Controladora estándar PCI IDE de doble canal	OK
IRQ 20	Controladora estándar PCI IDE de doble canal	OK
IRQ 21	Controladora de host universal USB de la familia Intel(R) ICH10 - 3A38	OK
IRQ 22	Controladora de High Definition Audio	OK
IRQ 23	Controladora de host mejorado USB de la familia Intel(R) ICH10 - 3A3A	OK
IRQ 23	Controladora de host universal USB de la familia Intel(R) ICH10 - 3A34	OK
IRQ 24	NVIDIA GeForce GTX 260	OK
IRQ 81	Sistema Microsoft compatible con ACPI	OK
IRQ 82	Sistema Microsoft compatible con ACPI	OK
IRQ 83	Sistema Microsoft compatible con ACPI	OK
IRQ 84	Sistema Microsoft compatible con ACPI	OK

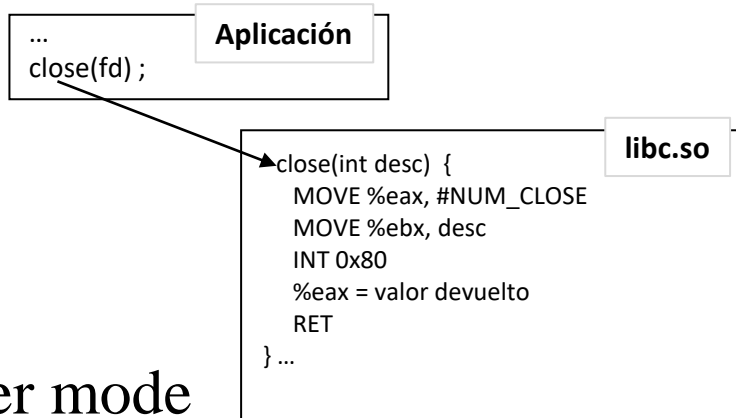
Interrupts by Software.

System calls and operating systems

- ▶ The system call mechanism is the one that allows user programs to request the services offered by the operating system
 - ▶ Load programs into memory for execution
 - ▶ Access to peripheral devices
 - ▶ Etc.
- ▶ Similar to the system calls offered by the CREATOR simulator
 - ▶ WepSIM examples show how system calls are internally implemented.

Interrupts by Software.

System calls (example: Linux)



User mode

Kernel mode

Interrupts by Software.

System calls (example: Linux)

...
close(fd);

Aplicación

close(int desc) {
 MOVE %eax, #NUM_CLOSE
 MOVE %ebx, desc
 INT 0x80
 %eax = valor devuelto
 RET
} ...

libc.so

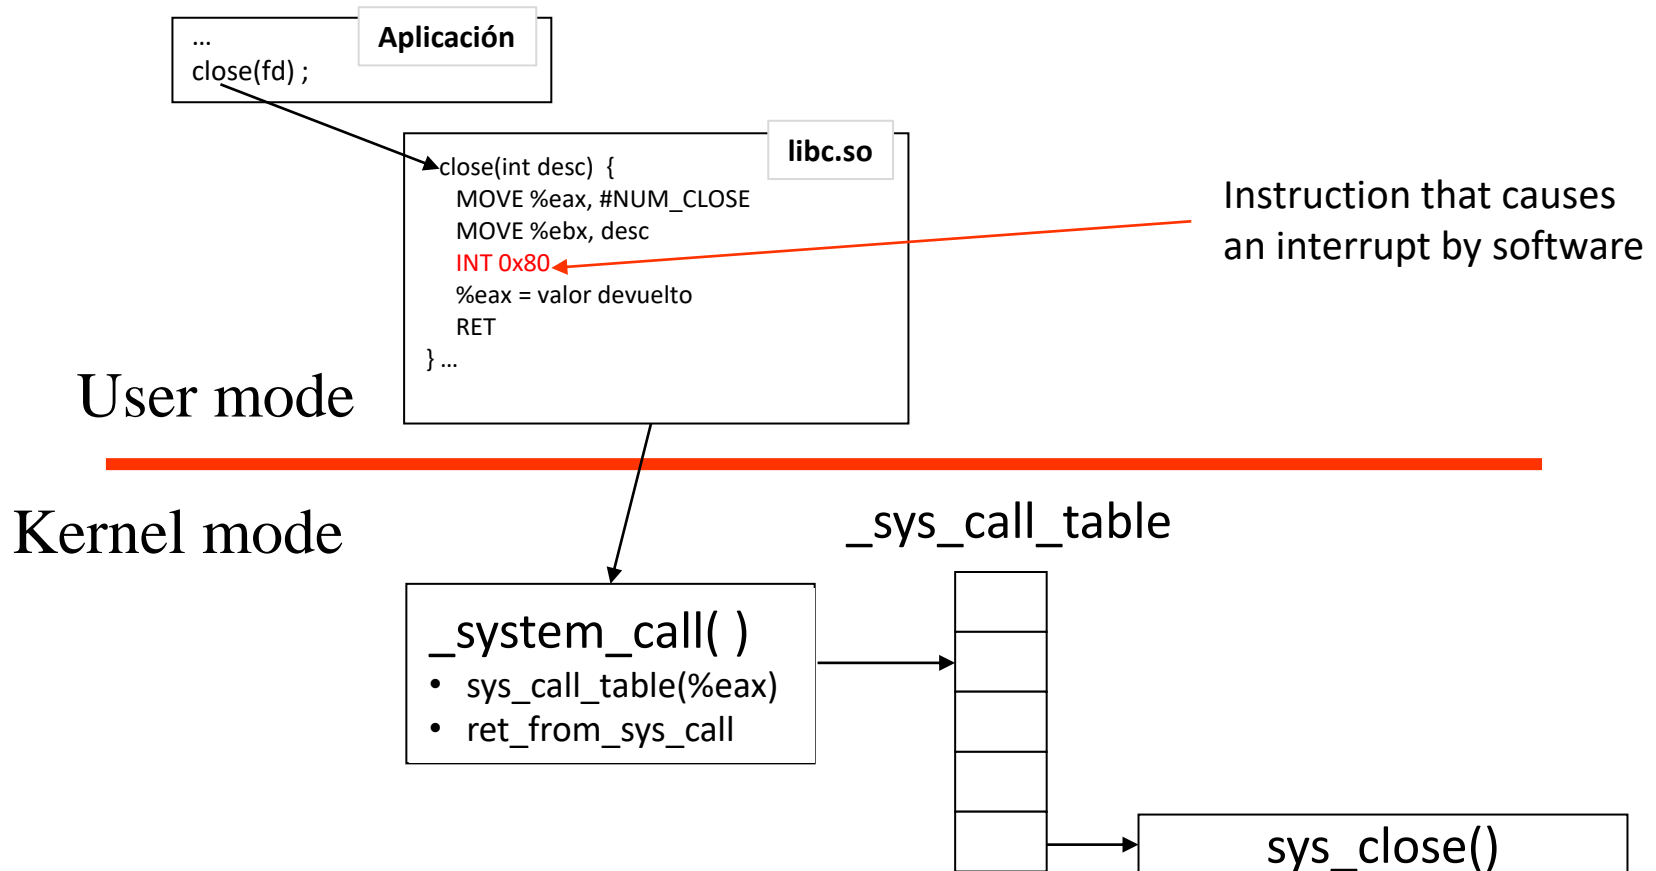
Instruction that causes
an interrupt by software

User mode

Kernel mode


Interrupts by Software.

System calls (example: Linux)



Interrupts by Software.

System calls (example: Linux)



Clock interrupts and operating system


- ▶ The signal that governs the execution of machine instructions is divided by a frequency divider to generate an external interruption every certain time interval (a few milliseconds)
- ▶ These **clock interruptions** or ticks are periodic interruptions that allow the operating system to come in and run periodically, preventing a user program from monopolizing the CPU
 - ▶ Allows to alternate the execution of various programs on a system given the appearance of simultaneous execution
 - ▶ Each time a clock interruption arrives, the program is suspended and the operating system that runs the scheduler is skipped to decide the next program to run

Contents

1. Computer elements
2. Processor organization
3. Control unit
4. Execution of instructions
5. Control unit design
6. Execution modes
7. Interrupts
8. Computer startup
9. Performance and parallelism


Computer booting

- ▶ The *Reset* loads the predefined values in registers:
 - ▶ $PC \leftarrow$ initial address of the **initialization program** (in ROM memory)



Computer booting

- ▶ The *Reset* loads the predefined values in registers:
 - ▶ PC ← initial address of the **initialization program** (in ROM memory)
- ▶ The **initialization program** is executed:
 - ▶ System test (POST)



```
Award Modular BIOS v6.00PG, An Energy Star Ally
Copyright (C) 1984-2007, Award Software, Inc.

Intel X38 BIOS for X38-DQ6 F4

Main Processor : Intel(R) Core(TM)2 Extreme CPU X9650 @ 4.00GHz(333x12)
<CPUID:0676 Patch ID:0000>
Memory Testing : 2096064K OK


Memory Runs at Dual Channel Interleaved
IDE Channel 0 Slave : WDC WD3200AAJS-00RYA0 12.01B01
IDE Channel 1 Slave : WDC WD3200AAJS-00RYA0 12.01B01

Detecting IDE drives ...
IDE Channel 4 Master : None
IDE Channel 4 Slave : None
IDE Channel 5 Master : None
IDE Channel 5 Slave : None

<DEL>:BIOS Setup <F9>:XpressRecoveryZ <F12>:Boot Menu <End>:Quit
09/19/2007-X38-ICH9-6A79060QC-00
```



Computer booting

- ▶ The *Reset* loads the predefined values in registers:
 - ▶ $PC \leftarrow$ initial address of the **initialization program** (in ROM memory)
- ▶ The **initialization program** is executed:
 - ▶ System test (POST)
 - ▶ Load into memory the **operating system loader** (MBR)




Computer booting

- ▶ The *Reset* loads the predefined values in registers:
 - ▶ PC ← initial address of the **initialization program** (in ROM memory)
- ▶ The **initialization program** is executed:
 - ▶ System test (POST)
 - ▶ Load into memory the **operating system loader** (MBR)
- ▶ The **Operating System Loader** is executed:
 - ▶ Sets boot options






Computer booting

- ▶ The *Reset* loads the predefined values in registers:
 - ▶ $PC \leftarrow$ initial address of the **initialization program** (in ROM memory)
- ▶ The **initialization program** is executed:
 - ▶ System test (POST)
 - ▶ Load into memory the **operating system loader** (MBR)
- ▶ The **Operating System Loader** is executed:
 - ▶ Sets boot options
 - ▶ Loads the **loading program**



Computer booting

- ▶ The *Reset* loads the predefined values in registers:
 - ▶ PC ← initial address of the **initialization program** (in ROM memory)
- ▶ The **initialization program** is executed:
 - ▶ System test (POST)
 - ▶ Load into memory the **operating system loader**
- ▶ The **Operating System Loader** is executed:
 - ▶ Sets boot options
 - ▶ Loads the **loading program**
- ▶ The **Loading Program** is executed:
 - ▶ Sets the initial state of the O.S.
 - ▶ Loads the O.S. and executed it.




```
Configuring ISA PNP
Setting system time from the hardware clock (localtime).
Using /etc/random-seed to initialize /dev/urandom.
Initializing basic system settings ...
Updating shared libraries
Setting hostname: enpc23.murdoch.edu.au
INIT: Entering runlevel: 4
rc.M ==> Going multiuser...
Starting system logger ... [ OK ]
Initialising advanced hardware
Setting up modules ... [ OK ]
Initialising network
Setting up localhost ... [ OK ]
Setting up inet1 ... [ OK ]
Setting up route ... [ OK ]
Setting up fancy console and GUI
Loading fc-cache ... [ OK ]
rc.v1init ==> Going to runlevel 4
Starting services of runlevel 4
Starting dnsmasq ... [ OK ]
==> rc.X Going to multiuser GUI mode ...
XFree86 Display Manager
Framebuffer /dev/fb0 is 307200 bytes.
Grabbing 640x480 ...
```

Computer booting

summary

- ▶ The *Reset* loads the predefined values in registers:
 - ▶ $PC \leftarrow$ initial address of the **initialization program** (in ROM memory)
- ▶ The **initialization program** is executed:
 - ▶ System test (POST)
 - ▶ Load into memory the **operating system loader** (MBR)
- ▶ The **Operating System Loader** is executed:
 - ▶ Sets boot options
 - ▶ Loads the **loading program**
- ▶ The **Loading Program** is executed:
 - ▶ Sets the initial state of the O.S.
 - ▶ Loads the O.S. and executed it.



Contents

1. Computer elements
2. Processor organization
3. Control unit
4. Execution of instructions
5. Control unit design
6. Execution modes
7. Interrupts
8. Computer startup
9. Performance and parallelism

Program execution time

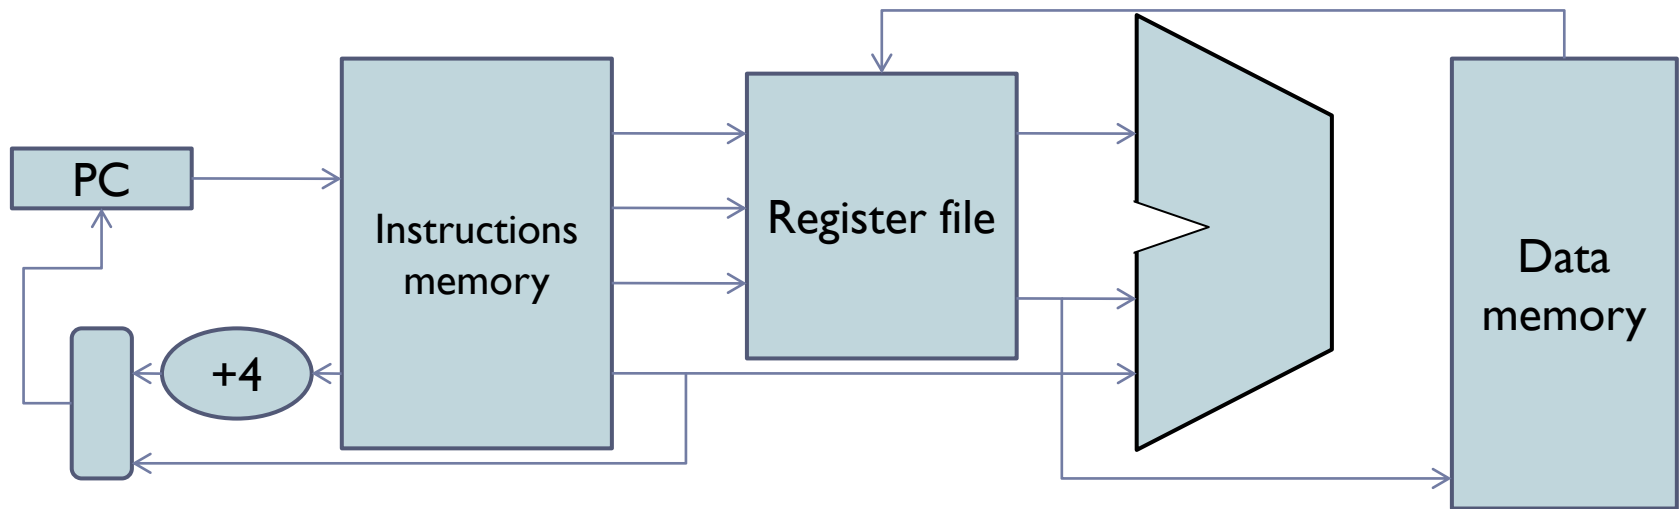
$$\text{Time}_{\text{execution}} = \text{IN} \times \text{CPI} \times t_{\text{cycle_CPU}} + \text{IN} \times \text{AMI} \times t_{\text{cycle_mem}}$$

- ▶ **IN** is the number of instructions of the program
- ▶ **CPI** is the average number of clock cycles to execute an instruction
- ▶ $t_{\text{cycle_CPI}}$ is the cycle clock duration
- ▶ **AMI** is the average number of memory access per instruction
- ▶ $t_{\text{cycle_mem}}$ is the time needed for a memory access

Factors affecting execution time

	IN	CPI	$t_{\text{cycle_CPI}}$	AMI	$t_{\text{cycle_mem}}$
Program	✓			✓	
Compiler	✓	✓		✓	
Instruction set	✓	✓	✓	✓	
Organization		✓	✓		✓
Technology			✓		✓

Model of processor based on datapath (without internal bus)



Instruction level parallelism

- ▶ Concurrent execution of several machine instructions
- ▶ Combination of elements working in parallel:
 - ▶ **Pipelined processor**: use pipelines in which multiple instructions are overlapped in execution
 - ▶ **Superscalar processor**: multiple independent instruction pipelines are used. Each pipeline can handle multiple instructions at a time
 - ▶ **Multicore processor**: several processors or cores in the same chip


Segmentation of instructions



- ▶ Stages in the execution of instructions:
 - ▶ **IF:** Instruction fetch
 - ▶ **D:** Decoding
 - ▶ **RO:** Read operands
 - ▶ **EX:** Execution
 - ▶ **WO:** Write operands

Segmentation of instructions


without pipeline



- ▶ Stages in the execution of instructions:
 - ▶ **IF:** Instruction fetch
 - ▶ **D:** Decoding
 - ▶ **RO:** Read operands
 - ▶ **EX:** Execution
 - ▶ **WO:** Write operands

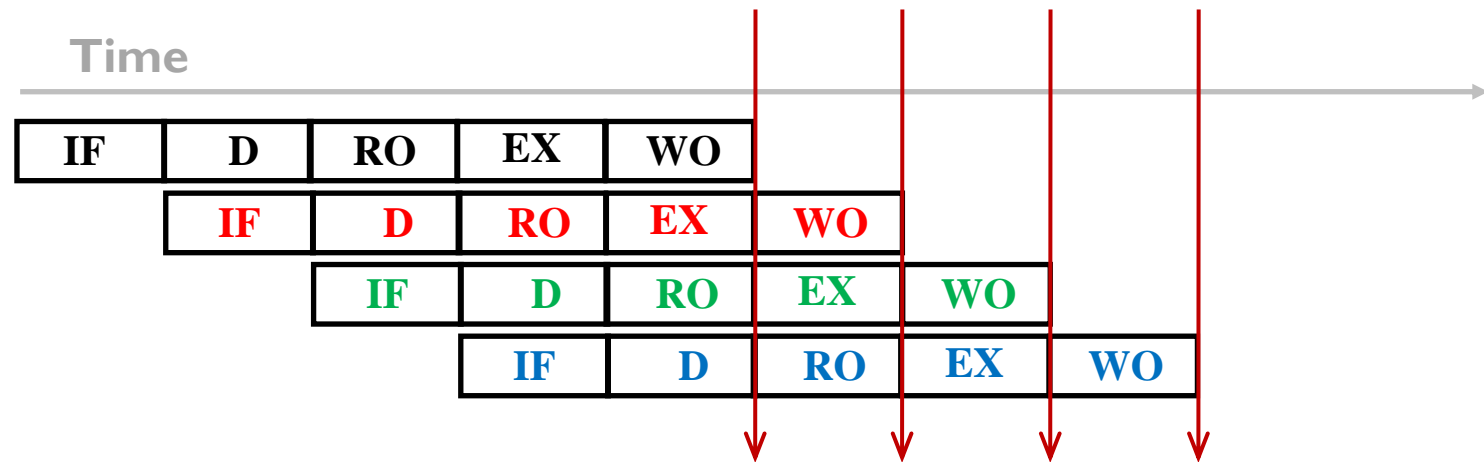
Segmentation of instructions

without pipeline




- ▶ If each phase takes N clock cycles, then:
 - ▶ One instruction takes $5 \cdot N$ clock cycles to be executed
 - ▶ $1/5$ of instruction is issued every N clock cycles

Segmentation of instructions with pipeline



- ▶ If each phase takes N clock cycles, then:
 - ▶ One instruction takes $5 \cdot N$ clock cycles to be executed
 - ▶ One instruction is issued every N clock cycles



Superscalar



- Pipeline with several functional units in parallel

Multicore

- Multiples processors in the same chip



Multicore

- ▶ Multiples processors in the same chip



3 Knights Landing Products

A Paradigm Shift for Highly-Parallel

	Intel® 64 / AVX-512	Intel® 64 / AVX-512	Intel® 64 / AVX-512
Programming Model	PCIe	Fabric	Integrated Fabric
I/O	Baseline	>25% Better¹	>25% Better¹
Power Efficiency	Baseline	Intel server-class	Intel server-class
Resiliency	>3 TF¹	>3 TF¹	>3 TF¹
Performance	up to 16GB	up to 400GB¹	up to 400GB¹
Memory Capacity	>5x STREAM vs. DDR4¹	>5x STREAM vs. DDR4¹	>5x STREAM vs. DDR4¹
Memory Bandwidth			

<http://wccfttech.com/intel-knights-landing-detailed-16-gb-highbandwidth-on-die-memory-384-gb-ddr4-system-memory-support-8-billion-transistors/>

ARCOS Group

uc3m | Universidad **Carlos III** de Madrid

Lesson 4 (III) The processor

Computer Structure
Bachelor in Computer Science and Engineering

